

Functions in PGSQL

functions in psql  
create or replace function funcName()  
return s returntype as \$\$  
begin  
raise notice 'Hello Birbl!';  
end;  
\$\$ language langName  
The two \$ are always necessary. Also note the  
returns with an s and the language at the end.  
Which MUST be a PROCEDURAL LANGUAGE,  
so c++ doesn't work here.

Parameters are handled like in any language  
func(x bigint, y bigint)  
you can also define multiple return types  
func(variadic a numeric[])  
or a generic return  
func(param anyelement)

Variable Declaration:  
returns void as \$\$  
DECLARE  
x bigint; y bigint;  
BEGIN ....

Variable manipulation: x := 6 + 4  
if: IF n = 0 THEN RETURN 1;  
(optional)ELSE RETURN 2; END IF;  
ELSIF also possible (note elsif not elseif)

case x when 1, 2 then msg := 'one or two'; end case;  
essentially this checks if x is 1 OR 2

case when x between 0 and 10 then ....  
similar but with a range, both can be  
simulated by if else.

Exceptions: BEGIN z:= x / y;  
EXCEPTION WHEN division-by-zero  
THEN z:= 0; (or error rather) END;  
if you want to catch all: WHEN others THEN  
often used after exception: RAISE; (show error)

For Loop: For var IN query LOOP  
statements END LOOP;  
for r in SELECT \* FROM ang LOOP  
RETURN NEXT r; END LOOP; RETURN; END;  
note that the return next doesn't return  
you store it in a buffer  
and return it at the end of the function.  
for infinite loops: FOR i IN 1...max LOOP;  
update and insert: INSERT INTO ANG VALUES(...);  
UPDATE ang set salary = salary + 500  
where name = 'dashie';  
interestingly, after the where name = 'dashie'  
you can use if not found then (handle error)  
this allows for easier error handling.

queries: execute 'SELECT \* from ang'  
---- into result; return result; END;  
comments are done by either -- or /\* \*/ for multiline

anonymous function: you can omit  
the name and just write do \$\$ ...

cursor: declare curs CURSOR FOR query;  
BEGIN OPEN curs; LOOP do something CLOSE  
curs; END;  
Cursors are essentially just iterables.  
cursors can also be unbound curs refcursor  
or they can be parameterized curs3 cursor(arg)

PL/pgSQL: Datentypen

• PL/pgSQL übernimmt alle SQL Datentypen:

- Boolean:
- Zahlen: int, integer, number
- Strings, Datum, etc.
- Arrays: alle Datentypen gefolgt von [], z.B. int[]
- Weitere: JSON, etc.

• ... ergänzt mit zusätzlichen Datentypen:

- var5 angestellter.idfntype; -- abgeleiteter col.-Typ
- var6 angestellter.frowtype; -- abgeleitet von Tabelle
- var7 record; -- generischer Record this is a simple entry
- var8 anyelement; -- generischer Typ gemäß Fn.-Argument, vgl. nachfolgend
- curs1 refcursor; curs2 cursor ...; -- vgl. nachfolgend

arrays: SELECT '1,2,3':int[]  
or SELECT ARRAY[1,2,3]  
var int[] only in variable declaration.  
!!arrays start with 1 in psql !!  
return types: all of the above AND void,  
SETOF type (array of a type), TABLE, Trigger

Arrays: Accessoren

create table tictactoe as  
(select i as id,  
array(  
[z1 k1, z1 k2, z1 k3],  
[z2 k1, z2 k2, z2 k3],  
[z3 k1, z3 k2, z3 k3])  
as board)  
);

Index Query: intuitiv wie eine Koordinate ("1-basiert": Start mit 1 nicht 0):  
select board[1][1] from tictactoe;  
-- z1 k1

Slice Query: "Untergrenze/Obergrenze" für jede Dimension:  
select board[2:3][1:1] from tictactoe;  
-- ((z2 k1), (z3 k1))

Max-Bound-Abkürzung "[2]" vermeiden (Verwechslungsgefahr), besser [1:2]  
select board[z1:2][2] from tictactoe; -- [2]>[1:2]  
-- ((z2 k1, z2 k2), (z3 k1, z3 k2))

Suche mit ANY:  
select \* from tictactoe where 'z2 k2' = any(board);  
-- 1: ((z1 k1, z1 k2), (z2 k1, z2 k2), (z3 k1, z3 k2))

Arrays: Operatoren

• «Is equal» =  
SELECT ARRAY[1,2,3] = ARRAY[1,2,3];  
true  
SELECT ARRAY[3,2,1] = ARRAY[1,2,3];  
false

«Contains» @>  
-- SELECT ARRAY[1,7,4,2,6] @> ARRAY[2,7];  
«Is contained by» @<  
-- SELECT ARRAY[2,7] @< ARRAY[1,7,4,2,6];  
«Overlaps» &&  
-- SELECT ARRAY[1,4,3] && ARRAY[2,1]

hstore / map:  
Create und Insert:  
CREATE TABLE test (id integer, col2 hstore, col3 text);  
  
INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL);  
  
SELECT \* FROM test;  
id| col2 | col3  
1| "a"=>"123", "b"=>"foo", "c"=>"bar"| null  
(1 row)  
  
Queries:  
- List all keys  
SELECT akeys(mykvpfld) FROM ...  
  
- Get all key-value pairs  
SELECT each(mykvpfld) FROM ...  
  
- Get key value (as text)  
SELECT mykvpfld->'name' FROM ...  
  
- Test if left hstore is contained in right hstore:  
... WHERE mykvpfld @> 'tourism=>zoo'; -- or hstore('tourism','zoo')

Operatoren:  
- ">" get value for key: SELECT 'a'=x, b=>y':hstore -> 'a'  
- @>, etc. ... ähnlich wie Array-Operator

hstore supports GIST/GIN indexing

FUNCTION

PROCEDURE

Use in an expression

Return a value

Return values as OUT parameters

Return a single result set

Return multiple result sets

Contain transactions

Make it run using ...

EXECUTE

CALL

also note that type->rowtype is used like this:  
r ang%rowtype -> for r in select \* from ang;

important to know, you can always use these  
functions to manipulate queries, for example  
select upper(name) from ang;  
depending on the function you can also  
select generatetab(1,10)

• Es gibt zusätzlich IN, OUT, INOUT  
create function foo(IN pl type)...

- IN: call by value; Variablen oder Ausdrücke als Argument
- OUT: call by reference; nur Variablen als Argument
- INOUT: beides

cast: cast(input as type);  
cast(record.id as text);

stored procedures are nothing but a chaining of functions:  
• Schritt 1 in PL/pgSQL: siehe Beispiel 3 (SP-Funktion mit in/out-Parametern)  
• Schritt 2: Deklaration in Java/JPA (aka Registrierung in JPA):  
@NamedStoredProcedureQuery(  
name = "MySum", -- JPA-Objekt  
procedureName = "mysum", -- Name der SP-Fn. (DB-Objekt)  
parameters = {  
@StoredProcedureParameter(mode = ParameterMode.IN, type = Double.class, name = "x"),  
@StoredProcedureParameter(mode = ParameterMode.IN, type = Double.class, name = "y"),  
@StoredProcedureParameter(mode = ParameterMode.OUT, type = Double.class, name = "sum")  
}  
)  
• Schritt 3: Call It!  
@NamedStoredProcedureQuery query = this.em.createNamedStoredProcedureQuery("MySum");  
query.setParameter("x", 1.23d);  
query.setParameter("y", 4.56d);  
query.execute();  
Double sum = (Double) query.getOutputParameterValue("sum");

some good to know things: plain SQL is more efficient.  
write variables lower case for sql  
use cast over typename -> not select date '2022-06-07'

Triggers

• sind DB-Objekte und immer einer Tabelle zugeordnet

• werden in Stored Procedures programmiert

• haben keine Parameter

• können nicht direkt aufgerufen werden

• werden vom DBMS beim Eintreten eines Events aufgerufen

• haben bei der Ausführung die Rechte ihres Owners

Trigger can pass parameters to function  
FOR EACH [statement | row]  
Events: INSERT, UPDATE, DELETE, TRUNCATE, INSTEAD OF  
Function executes BEFORE or AFTER changes

statement is once, row means once  
per row, aka for the entire table.

• Before triggers can change  
contents of new row

• After triggers can only  
respond to what has  
happened  
e.g. Foreign Keys

• Return values of AFTER  
triggers will be ignored

• Triggers execute in  
alphabetical order

The INSTEAD Trigger can be use  
to avoid crashes:  
-> INSTEAD OF UPDATE  
-> ON UPDATE DO INSTEAD  
for example, trying to update  
a read-only view -> INSTEAD OF  
UPDATE

DDL für Triggers  
CREATE | DROP | ALTER TRIGGER  
mytrigger ...  
ON mytable ...

PL/pgSQL: Trigger-Fn.-Variablen

• TG\_NAME Name des Triggers (TG)  
• TG\_WHEN BEFORE oder AFTER  
• TG\_LEVEL ROW od. STATEMENT  
• TG\_OP INSERT, UPDATE, DELETE, (TRUNCATE)  
• TG\_REID Old der Tabelle  
• TG\_RELNAME Name der Tabelle  
• TG\_TABLE\_SCHEMA Schema der Tabelle

CREATE OR REPLACE FUNCTION dt\_trigger\_func()  
RETURNS TRIGGER AS \$\$  
BEGIN  
IF (TG\_OP = 'INSERT') THEN  
NEW.creation\_date := now();  
ELSEIF (TG\_OP = 'UPDATE') THEN  
NEW.modification\_date := now();  
END IF;  
RETURN NEW;  
END  
\$\$ LANGUAGE plpgsql;  
CREATE TRIGGER dt\_trigger  
BEFORE INSERT OR UPDATE  
ON mytable  
FOR EACH ROW  
EXECUTE PROCEDURE dt\_trigger\_func();

Syntax-Beispiel Trigger-Fn. passend zu Trigger:  
CREATE FUNCTION mytriggerfn()  
RETURNS TRIGGER --  
AS \$\$  
<body>  
\$\$ language plpgsql;  
-- Eine Trigger-Fn. hat keine Fn.-Parameter  
-- Diese werden über Trigger-Fn.-Variablen übergeben, u.a.:

- TG\_NARGS Anzahl Parameter
- TG\_ARGV[] Array von Parametern als TEXT

Return types: RETURN NEW -> returns a new table/row  
RETURN OLD -> returns the old table/row  
(but could change other rows!)  
RETURN NULL -> cancel operation.

running order: before statement, before row,  
after row, after statement ->  
and of course alphabetically.

inside the Trigger functions you can use the  
variables that don't matter aka can be ANY -> user  
or the entered user from the trigger -> NEW.user  
or explicitly the old one -> OLD.user  
and last user defined stuff like -> SELECT 'I'  
which just places an I as the variable  
or something like now() for timestamps.

Triggers make the database slower and harder to maintain.  
some databases therefore let you disable them if you want.  
On a table basis.  
Also watch out for cascading effects of triggers  
they might cause something else to be deleted.

Stored Procedures are really helpful for security  
They have all the privileges, but only  
allow the user to do what the  
creator has predefined.

User defined types in SQL

CREATE DOMAIN ...

CREATE TYPE ...

• erstellt einen benutzerdefinierten  
Datentyp

• Für Datenschema und Stored  
Procedures

• mit Constraints wie NOT NULL, CHECK

create domain contact\_name as  
varchar(255) not null  
check (value ~ '^s');

create type eastereggs as (  
outer: text,  
inner: text,  
create type traffic\_light\_t as  
enum('red', 'yellow', 'green');

One is a completely new type, the other is just a type that is  
made up of already known types.  
Essentially the right one is a class of struct.  
The left is a completely new thing, that is not yet implemented.

Optimization and indexing

The basics of indexing is that it saves  
time on queries, but it uses more space,  
and needs to be redone on update/insert.

Data stored in Pages / Heap (Collection of Pages)  
Indexing on either can be possible.  
ex: Page index -> Primary key index  
ex: Heap Index -> index on tables  
PSQL does table cluster indexing  
instead of integrated indexing -> key value (both indexed)

B-Tree: the default, can index multiple entries (only btree!)  
CREATE UNIQUE INDEX name ON table (column [1,2,...])

Hash Index: just like hashmap in programming  
good for single or small multiple queries  
bad for entire tables etc -> collision  
hashing might take a long time with a lot of data.  
B-tree almost always better!

GIST: balanced/treelike , Range/neighbor/fulltext search  
used for geometric datatypes. SP-Gist for unbalanced trees.

GIN: General Inverted Index "list of words  
that point to documents" wtf?. Good for duplicates.  
Good for hstore,Json,Arrays  
Vergleich zu GIST:

- Ca. 3x schnellerer Zugriff
- Ca. 2 - 3x mehr Diskplatz
- Ca. 2 - 3x länger bis Index erstellt ist
- Ca. 10x langsamer bei Update

Bitmap indexing: Bitmap -> 0 1 stores Booleans/Enums  
very fast read / slow update  
in postgres only implicit use Brin instead.

BRIN: Block Range Index, stores min/max values as blocks  
good for range search, sorted data, small disk usage  
data is naturally sorted, address next to postal code.

Bloom Index -> equality search,  
Trigram Index -> Full text search  
RUM , non-default-GIN jsonb-path-ops

creating index:  
CREATE INDEX <indexname> ON <table(attribute)>;  
and: DROP INDEX <indexname>;  
default index order for psql is btree,ASC,NULL first

Index-Variationen

• Zusammengesetzter Index  
- Bsp.: CREATE INDEX idx\_addr ON addr(phone,name); kann für Queries auf phone  
und "phone AND name" genutzt werden, sowie für bestimmte Q. auf name;  
jedoch nicht für Suffix-Q. "... LIKE %name";

• Index mit INCLUDE  
- Bsp.: CREATE INDEX idx\_addr ON addr(phone) INCLUDE (name);

Index-Variationen ff.

• Partieller Index  
- Bsp.: CREATE INDEX idx\_addr ON addr(status) WHERE status='active';  
index on function()

• Indexe mit Funktionen / Ausdrücke ("Funktionaler I.")  
- Bsp.: CREATE INDEX idx\_addr ON addr(lower(name));

• Nicht nur PostgreSQL

PG planner join strategies: Nested Loop,Merge,Hash  
-Nested Loop: for r in right row r == for l in left row....  
good for small tables, easy to setup  
-Merge: Merge rows one after the other  
higher starting cost, good for bigger tables  
-Hash: Hash the row then compare to other row  
only == check possible, high starting cost, low execution cost

PG planner scans: Full,Index,Index Only, Bitmap  
-full scans the entire table  
-index scans index and more (if necessary)  
-index only only scans index  
-Bitmap scans the bitmap generated by an index.

The steps of optimization  
1. generate the plan of transaction  
2. reform the term to optimize performance  
without knowledge of the internal structure.  
-> all values are considered equal  
3. optimization based on: available indexes, analysis and costs  
4. generate all possible plans to calculate cost  
5. analyze said plans -> how many tuples, what kind...  
6. profit?

selectivity this is the ratio of tuples a query returns  
low selectivity would mean high number of rows  
an example is select \* from table where sex='Male'  
the opposite would be high selectivity.

density this is the ratio of duplicates a query returns  
the more duplicates the higher the density.  
you can therefore also make graphs about the distribution  
of density -> names a-c low g-l high

best practices: index only when selectivity over 10%  
numeric comparison over text, join over subquery,  
use short attributes, understand the query -> user,  
don't select \* from, don't use cross products.

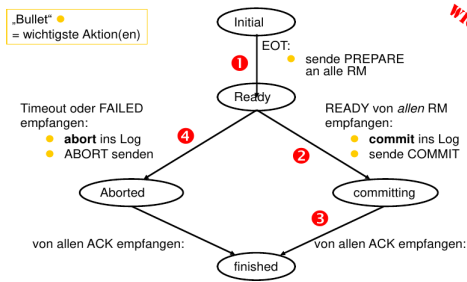
distributed database systems  
the advantages and disadvantages are obvious:  
+ better performance - more complex  
+ better reliability + better management

homogeneous database system:  
> all nodes have identical software  
> all nodes know about each other and work together  
> shows itself to the user as one big system

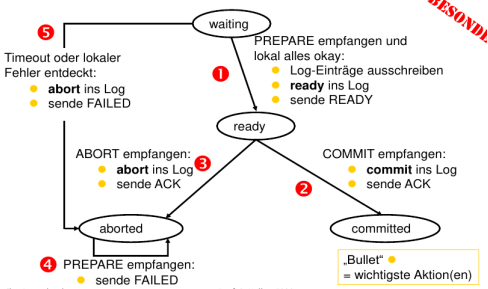
heterogeneous database system:  
> nodes can have different software  
> as well as different schemas -> can lead to problems  
> nodes might NOT know about each other.

<b>fragmentation:</b> this is the splitting of schemas into multiple Nodes -> table 1 in node1 table 2 in node 2.	
In Psql the horizontal fragmentation happens in 3 ways: > 1 or more attributes for "partitioning key" > "list" explicit designation > hash function (ex: Modulo)	
In Graph stores this would be called "sharding" an example for this is the MongoDB horizontal partitioning and allocation within a single node	
<b>replication:</b> this is the duplication of data in schemas this means table 1 might be on node1 and 2. <b>vertical -&gt; splitting of columns</b> row1 in node1, row2 in node2 <b>horizontal -&gt; splitting of rows</b> part of column in node1 part of column in node2	
<b>unidirectional:</b>	
<b>bidirectional:</b>	
<b>synchronous:</b>	
<b>asynchronous:</b>	
<b>allocation:</b> this is the distribution of work to the nodes node1 might handle query or part of query1 while node2 handles something else.	
as already stated above, the user doesn't see anything about fragmentation or similar, the user simply interacts with the schema and executes transactions. these <b>transactions are always local</b> This means that the user will talk to <b>Transaction Manager</b> that will handle the transaction and call the necessary functions on the <b>Resource Managers</b> (inside nodes).	

Zustandsübergänge 2PC: Transaction Manager (TM)



Zustandsübergänge 2PC ff.: Resource Manager (RM)



<b>case TM failed/restart:</b> > if the TM crashes before the commit message -> abort > if the TM crashes after RM respond ready -> block RM this is one of the main problems btw...	
<b>case RM failed/restart:</b> > if no entry in log, RM aborts > if READY-Entry available -> RM asks TM what to do. > if COMMIT-Entry available -> RM redoes transaction	
<b>case message dropped:</b> > if the prepare statement gets lost, or the RM doesn't respond, then the TM simply aborts the transaction for all. > if RM doesn't get a response in READY state then the RM will remind the TM until it gets one.	



	<p>easiest way to convert PSQL to JSON is a temporary table.</p> <pre>create temporary table angprojj as select   ang.persnr,   jsonb_build_object(     'persnr', ang.persnr,     'name', min(ang.name),     'projects', jsonb_agg(trim(proj.bezeichnung))   ) as angwithproj</pre> <p>joins are omitted, would be below...</p> <p>you can also directly <code>select json_build_object</code> or <code>select json_build_object(jsonb_agg(tmp))</code> also note the <code>jsonb_agg(trim (proj.bezeichnung))</code> for simplicity, the joins on the picture are removed.</p> <p>The jsonb agg is necessary for aggregation otherwise the join wouldn't work for JSON.</p>
	<p><b>JSON-Daten abfragen</b></p> <pre>select persnr, angwithproj from angprojj where angwithproj-&gt;&gt;'persnr' = 1001::text;  select persnr, angwithproj from angprojj where (angwithproj-&gt;&gt;'name') like 'Marxer%'  select persnr, jsonb_pretty(angwithproj) from angprojj where angwithproj-&gt;'projects' @&gt;   to_jsonb('Uranus'::text)</pre> <p>Get as text:</p> <pre>select persnr, angwithproj-&gt;'projects' as projects from angprojj; =&gt; 23 rows</pre> <p>CROSS JOIN zweier Tabellen. <code>jsonb_array_elements_text()</code> gibt „setof text“ zurück:</p> <pre>select persnr, angwithproj-&gt;&gt;'name' as persname, value as projname from angprojj, jsonb_array_elements_text(angwithproj-&gt;'projects'); =&gt; 29 rows</pre> <p>JOIN LATERAL = CROSS JOIN und Boolean: Output identisch mit oben :</p> <pre>select persnr, angwithproj-&gt;&gt;'name' as persname, value as projname from angprojj cross join lateral jsonb_array_elements_text(angwithproj-&gt;'projects'); =&gt; 29 rows</pre>
	<p>JSON — SQL</p>