

Functions in PGPSQL
functions in psql create or replace function funcName() return s returntype as \$\$ begin raise notice 'Hello Birbl!'; end; \$\$ language langName The two \$ are always necessary. Also note the returns with an s and the language at the end. Which MUST be a PROCEDURAL LANGUAGE , so c++ doesn't work here.
Parameters are handled like in any language func(x bigint, y bigint) you can also define multiple return types func(variadic a numeric[]) or a generic return func(param anyelement)
Variable Declaration: returns void as \$\$ DECLARE x bigint; y bigint; BEGIN
Variable manipulation: x := 6 + 4 if: IF n = 0 THEN RETURN 1; (optional)ELSE RETURN 2; END IF; ELSIF also possible (note elsif not elseif)
case x when 1, 2 then msg := 'one or two'; end case; essentially this checks if x is 1 OR 2
case when x between 0 and 10 then similar but with a range, both can be simulated by if else.
Exceptions: BEGIN z:= x / y; EXCEPTION WHEN division-by-zero THEN z:= 0; (or error rather) END; if you want to catch all: WHEN others THEN often used after exception: RAISE; (show error)
For Loop: For var IN query LOOP statements END LOOP; for r in SELECT * FROM ang LOOP RETURN NEXT r; END LOOP; RETURN; END; note that the return next doesn't return you store it in a buffer and return it at the end of the function. for infinite loops: FOR i IN 1..max LOOP; update and insert: INSERT INTO ANG VALUES(...); UPDATE ang set salary = salary + 500 where name = 'dashie'; interestingly, after the where name = 'dashie' you can use if not found then (handle error) this allows for easier error handling.
queries: execute 'SELECT * from ang' —— into result; return result; END; comments are done by either -- or /* */ for multiline
anonymous function: you can omit the name and just write do \$\$...
cursor: declare curs CURSOR FOR query; BEGIN OPEN curs; LOOP do something CLOSE curs; END; Cursors are essentially just iterables. cursors can also be unbound curs1 refcursor or they can be parameterized curs3 cursor(arg)
<ul style="list-style-type: none">PL/pgSQL: Datentypen<ul style="list-style-type: none">Boolean:Zahlen: int, integer, numberStrings, Datum, etc.Arrays: alle Datentypen gefolgt von „[]“, z.B. int[]Weitere: JSON, etc.... ergänzt mit zusätzlichen Datentypen:<ul style="list-style-type: none">var5 angestellter.id%type; -- abgeleiteter col-Typvar6 angestellter.krowtype; -- abgeleitet von Tabellevar7 record; -- generischer Record This is a simple entry inside a query aka 1 personvar8 anyelement; -- generischer Typ gemäss Fn-Argument, vgl. nachfolgendcurs1 refcursor; curs2 cursor ...; -- vgl. nachfolgend arrays: SELECT '1,2,3':int[] or SELECT ARRAY[1,2,3] var int[] only in variable declaration. !!arrays start with 1 in psql !! return types: all of the above AND void, SETOF type (array of a type), TABLE, Trigger
Arrays: Accessoren <pre>create table tictactoe as (select s as lv (['23 k1', '23 k2'], ['23 k1', '23 k2'], ['23 k1', '23 k2'])) as board);</pre> Index Query: intuitiv wie eine Koordinate ("1-basiert": Start mit 1 nicht 0): select board[1][1] from tictactoe; -- (2 k1, 23 k1) Slice Query: "Untergrenze/Obergrenze" für jede Dimension: select board[2:3][1:1] from tictactoe; -- ((2 k1), (23 k1)) Max-Bound-Abkürzung "[2]" vermeiden (Verwechslungsgefahr), besser [1:2] select board[2:3][2] from tictactoe; -- [2]>=>[1:2] -- ((2 k1, 23 k2), (23 k1, 23 k2)) Suche mit ANY: select * from tictactoe where '2 k2' = any(board); -- 1:((21 k1, 21 k2 k1), (2 k1, 23 k2), (23 k1, 23 k2))
Arrays: Operatoren <ul style="list-style-type: none">• «is equal» =<ul style="list-style-type: none">– SELECT ARRAY[1,7,4,2,6] @> ARRAY[2,7];– SELECT ARRAY[1,2,3] = ARRAY[1,2,3];– true– SELECT ARRAY[3,2,1] = ARRAY[1,2,3];– false• «is contained by» <@<ul style="list-style-type: none">– SELECT ARRAY[1,7,4,2,6] <@ ARRAY[1,7,4,2,6];– true– SELECT ARRAY[1,2,3] <@ ARRAY[1,7,4,2,6];– false• «Overlap» &&<ul style="list-style-type: none">– SELECT ARRAY[1,4,3] && ARRAY[2,1]
hstore / map: Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 1 "a">=>"123", "b">=>"foo", "c">=>"bar" null (1 row) Queries: – List all keys SELECT akeys(mykvpfld) FROM ... – Get all key-value pairs SELECT each(mykvpfld) FROM ... – Get key value (as text) SELECT mykvpfld->'name' FROM ... – Test if left hstore is contained in right hstore: ... WHERE mykvpfld @> 'tourism=>200'; -- or hstore('tourism','200')
Operatoren: – ">" get value for key : SELECT 'a=x, b=y':hstore -> 'a' – "@>", etc. ... ähnlich wie Array-Operator hstore supports GIST/GIN indexing

	FUNCTION	PROCEDURE
Use in an expression	✓	✗
Return a value	✓	✗
Return values as OUT parameters	✓ (PG Spezialität)	✓ (PG v14)
Return a single result set	✓ (as table fn.)	✓
Return multiple result sets	✗	✓
Contain transactions	✗	✓
Make it run using ...	EXECUTE	CALL

also note that type%rowtype is used like this:
r ang%rowtype -> for r in select * from ang;
important to know, you can always use these
functions to manipulate queries, for example
select upper(name) from ang;
depending on the function you can also
select generatetab(1,10)

- Es gibt zusätzlich IN, OUT, INOUT
create function foo(IN pl type)...
 - IN: call by value; Variablen oder Ausdrücke als Argument
 - OUT: call by reference; nur Variablen als Argument
 - INOUT: beides

cast: cast(input as type);
cast(record.id as text);

stored procedures are nothing but a chaining of functions:
• Schritt 1 in PL/pgSQL: siehe Beispiel 3 (SP-Funktion mit in/out-Parametern)
• Schritt 2: Deklaration in Java/JPA (aka Registrierung in JPA):
@NamedStoredProcedureQuery(
name = "MySum", -- JPA-Objekt
procedureName = "mysum", -- Name der SP-Fn. (DB-Objekt)
parameters = {
 @StoredProcedureParameter(mode = ParameterMode.IN, type = Double.class, name = "x"),
 @StoredProcedureParameter(mode = ParameterMode.IN, type = Double.class, name = "y"),
 @StoredProcedureParameter(mode = ParameterMode.OUT, type = Double.class, name = "sum")
}
)
• Schritt 3: Call It!
@NamedStoredProcedureQuery query = this.em.createNamedStoredProcedureQuery("MySum");
query.setParameter("x", 1.23d);
query.setParameter("y", 4.56d);
query.execute();
Double sum = (Double) query.getOutputParameterValue("sum");

some good to know things: plain SQL is more efficient.
write variables lower case for sql
use cast over typename -> not select date '2022-06-07'

| Triggers |
| - sind DB-Objekte und immer einer Tabelle zugeordnet - werden in Stored Procedures programmiert - haben keine Parameter - können nicht direkt aufgerufen werden - werden vom DBMS beim Eintreten eines Events aufgerufen - haben bei der Ausführung die Rechte ihres Owners |
| - Trigger can pass parameters to function - FOR EACH [statement | row] - Events: INSERT, UPDATE, DELETE, TRUNCATE, INSTEAD OF - Function executes BEFORE or AFTER changes statement is once, row means once per row, aka for the entire table. |
| - Before triggers can change contents of new row - After triggers can only respond to what has happened e.g. Foreign Keys - Return values of AFTER triggers will be ignored - Triggers execute in alphabetical order The INSTEAD Trigger can be used to avoid crashes: -> INSTEAD OF UPDATE -> ON UPDATE DO INSTEAD for example, trying to update a read-only view -> INSTEAD OF UPDATE |
| - DDL für Triggers CREATE | DROP | ALTER TRIGGER mytrigger ... ON mytable ... FOR EACH ROW EXECUTE PROCEDURE mytriggerfn(); **PL/pgSQL: Trigger-Fn.-Variablen** - TG_NAME Name des Triggers (TG) - TG_WHEN BEFORE oder AFTER - TG_LEVEL ROW od. STATEMENT - TG_OP INSERT, UPDATE, DELETE, (TRUNCATE) - TG_REID OID der Tabelle - TG_RELNAME Name der Tabelle - TG_TABLE_SCHEMA Schema der Tabelle CREATE OR REPLACE FUNCTION dt_trigger_func() RETURNS TRIGGER AS \$\$ BEGIN IF (TG_OP = 'INSERT') THEN NEW.creation_date := now(); ELSEIF (TG_OP = 'UPDATE') THEN NEW.modification_date := now(); END IF; RETURN NEW; END \$\$ LANGUAGE plpgsql; CREATE TRIGGER dt_trigger BEFORE INSERT OR UPDATE ON mytable FOR EACH ROW EXECUTE PROCEDURE dt_trigger_func(); |
| **Syntax-Beispiel Trigger-Fn. passend zu Trigger:** CREATE FUNCTION mytriggerfn() RETURNS TRIGGER -- AS \$\$ <body> \$\$ language plpgsql; – Eine Trigger-Fn. hat keine Fn.-Parameter – Diese werden über Trigger-Fn.-Variablen übergeben, u.a.: - TG_NARGS Anzahl Parameter - TG_ARGV[] Array von Parametern als TEXT |
| Return types: RETURN NEW -> returns a new table/row RETURN OLD -> returns the old table/row (but could change other rows!) RETURN NULL -> cancel operation. |
| running order: before statement, before row, after row, after statement -> and of course alphabetically. |
| inside the Trigger functions you can use the variables that don't matter aka can be ANY -> user or the entered user from the trigger -> NEW.user or explicitly the old one -> OLD.user and last user defined stuff like -> SELECT 'I' which just places an I as the variable or something like now() for timestamps. |
| Triggers make the database slower and harder to maintain. some databases therefore let you disable them if you want. On a table basis. Also watch out for cascading effects of triggers they might cause something else to be deleted. |
| Stored Procedures are really helpful for security They have all the privileges, but only allow the user to do what the creator has predefined. |

User defined types in SQL
CREATE DOMAIN ... <ul style="list-style-type: none">• erstellt einen benutzerdefinierten Datentyp• Für Datenschema und Stored Procedures• mit Constraints wie NOT NULL, CHECK create domain contact_name as varchar(255) not null check (value != '\s');
CREATE TYPE ... <ul style="list-style-type: none">• erstellt einen zusammengesetzten Datentyp• Für Datenschema und Stored Procedures• Auch für ENUM Type verwendet• Keine Constraints-Angaben create type eastereggs as (outer: text, inner: text, create type traffic_light_t as enum('red', 'yellow', 'green');
One is a completely new type, the other is just a type that is made up of already known types. Essentially the right one is a class or struct. The left is a completely new thing, that is not yet implemented.
Optimization and indexing
The basics of indexing is that it saves time on queries, but it uses more space, and needs to be redone on update/insert.
Data stored in Pages / Heap (Collection of Pages) Indexing on either can be possible. ex: Page index -> Primary key index ex: Heap Index -> index on tables PSQL does table cluster indexing instead of integrated indexing -> key value (both indexed)
B-Tree: the default, can index multiple entries (only btree!) CREATE UNIQUE INDEX name ON table (column [1,2,...])
Hash Index: just like hashmap in programming good for single or small multiple queries bad for entire tables etc -> collision hashing might take a long time with a lot of data. B-tree almost always better!
GiST: balanced/treelike , Range/neighbor/fulltext search used for geometric datatypes. SP-Gist for unbalanced trees.
GIN: General Inverted Index "list of words that point to documents" wtf?. Good for duplicates. Good for hstore,Json,Arrays Vergleich zu GiST: <ul style="list-style-type: none">– Ca. 3x schnellerer Zugriff– Ca. 2 – 3x mehr Diskplatz– Ca. 2 – 3x länger bis Index erstellt ist– Ca. 10x langsamer bei Update
Bitmap indexing: Bitmap -> 0 1 stores Booleans/Enums very fast read / slow update in postgres only implicit use Brin instead.
BRIN: Block Range Index, stores min/max values as blocks good for range search, sorted data, small disk usage data is naturally sorted, address next to postal code.
Bloom Index -> equality search Trigram Index -> Full text search RUM , non-default-GIN jsonb-path-ops
creating index: CREATE INDEX <indexname> ON <table(attribute)>; and: DROP INDEX <indexname>; default index order for psql is btree,ASC,NULL first
Index-Variationen <ul style="list-style-type: none">• Zusammengesetzter Index – Bsp.: CREATE INDEX idx_addr ON addr(phone,name); kann für Queries auf phone und "phone AND name" genutzt werden, sowie für bestimmte Q. auf name; jedoch nicht für Suffix-Q. "... LIKE %name;"• Index mit INCLUDE Similar – Bsp.: CREATE INDEX idx_addr ON addr(phone) INCLUDE (name);
Index-Variationen ff. <ul style="list-style-type: none">• Partieller Index index if(condition) – Bsp.: CREATE INDEX idx_addr ON addr(status) WHERE status='active'; index on function()• Indexte mit Funktionen / Ausdrücke ("Funktionaler I.") – Bsp.: CREATE INDEX idx_addr ON addr(lower(name));• Nicht nur PostgresSQL!
PG planner join strategies: Nested Loop,Merge,Hash - Nested Loop: for r in right row r == for l in left row.... good for small tables, easy to setup - Merge: Merge rows one after the other higher starting cost, good for bigger tables - Hash: Hash the row then compare to other row equality check only , high starting cost, low execution cost
PG planner scans: Full,Index,Index Only, Bitmap - full scans the entire table - index scans index and more (if necessary) - index only scans index - Bitmap scans the bitmap generated by an index.
The steps of optimization
1. generate the plan of transaction 2. reform the term to optimize performance without knowledge of the internal structure. -> all values are considered equal 3. optimization based on: available indexes, analysis costs 4. generate all possible plans to calculate cost 5. analyze said plans -> how many tuples, what kind... 6. profit?
selectivity this is the ratio of tuples a query returns low selectivity would mean high number of rows an example is select * from table where sex='Male' the opposite would be high selectivity.
density this is the ratio of duplicates a query returns the more duplicates the higher the density. you can therefore also make graphs about the distribution of density -> names a-c low g-l high
best practices: index only when selectivity over 10% numeric comparison over text, join over subquery, use short attributes, understand the query -> user, don't select * from, don't use cross products.
distributed database systems
the advantages and disadvantages are obvious: + better performance – more complex +better reliability + better management
homogeneous database system: > all nodes have identical software > all nodes know about each other and work together > shows itself to the user as one big system
heterogeneous database system: > nodes can have different software > as well as different schemas -> can lead to problems > nodes might NOT know about each other.
heterogeneous systems can be federated or unfederated > unfederated -> no local users > federated -> either tightly coupled with global schema or loosely coupled using export schema

fragmentation: this is the splitting of schemas into multiple Nodes -> table 1 in node1
table 2 in node 2.

In Psql the horizontal fragmentation happens in 3 ways:
> 1 or more attributes for "partitioning key"
> "list" explicit designation
> hash function (ex: Modulo)

In Graph stores this would be called "sharding"
an example for this is the MongoDB
horizontal partitioning and allocation within a single node

replication: this is the duplication of data in schemas
this means table 1 might be on node1 and 2.

vertical -i splitting of columns
row1 in node1, row2 in node2

horizontal -i splitting of rows
part of column in node1 part of column in node2

unidirectional: Single-Master

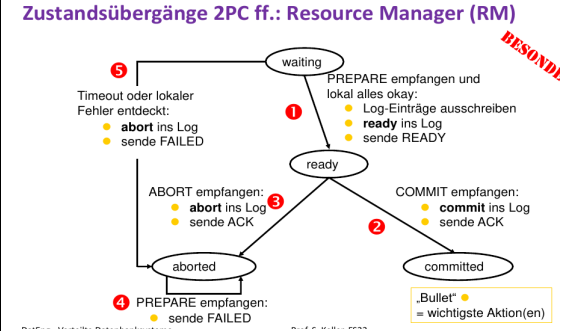
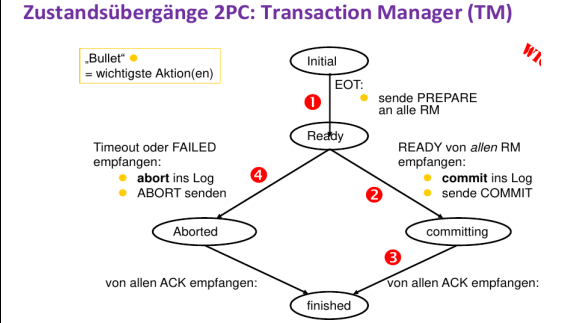
bidirectional: Multi-Master

synchronous OR asynchronous

allocation: this is the distribution of work to the nodes
node1 might handle query or part of query1
while node2 handles something else.

as already stated above, the user doesn't see anything about fragmentation or similar, the user simply interacts with the schema and executes transactions.
these **transactions are always local**
This means that the user will talk to **Transaction Manager** that will handle the transaction and call the necessary functions on the **Resource Managers** (inside nodes).

Two Phase Commit Protocol



case TM failed/restart:
> if the TM crashes before the commit message -> abort
> if the TM crashes after RM respond ready -> block RM
this is one of the main problems btw...

case RM failed/restart:
> if no entry in log, RM aborts
> if READY-Entry available -> RM asks TM what to do.
> if COMMIT-Entry available -> RM redoes transaction

case message dropped:
> if the prepare statement gets lost, or the RM doesn't respond, then the TM simply aborts the transaction for all.
> if RM doesn't get a response in READY state then the RM will remind the TM until it gets one.

There are different 2PC protocols, PSQL, MYSQL, X/Open, Java Transaction API, Oracle Transaction Manager, Microtroll...
+ often used / proven system + guarantees ACID
- slow - blocks transactions often
!!only use this when the complexity calls for it!!

Query handling:
there are 3 ways of handling a queries distributed databases:
> Ship Whole -> run query locally and remove duplicates
> Push Down -> split query when needed
> Fetch-as-needed -> Send join attributes to the correct node

NoSQL: Key/Value Stores

Why even NoSQL? -> fast, lots of data, webbased, scalable
RDBMS -> relational has the problem of complexity
it doesn't integrate seamlessly into programming languages
it therefore doesn't make sense to use this as something like a webstore. -> !!! OR MAPPING !!!

attributes of NoSQL:
> easy to use API (http) > made for big volumes of data
> provided as cloud storage > not relational, schema free
> BASE instead of ACID > no query norm (other than SQL)

The Cap Theorem
The CAP theorem dicates, that you can only have 2 of 3 desired traits of a database, these include **consistency**, **partition tolerance** and **availability**.
traditional databases have the **ACID** philosophy which is both **consistent** and **tolerant**, however it often blocks transactions.
here comes **NoSQL with the BASE** philosophy
it is both **available** at all times and **tolerant**
however it is **NOT consistent during transactions**.
Only after those have stopped will the system become consistent.
>tolerance means the system will work despite partial outage<

BASE THEOREM
>> **Basically Available:**
The database will not block transactions
any and all requests will be responded to (can still fail though!)
> **Soft state:**
The state of the system can change even without input, therefore we consider this to be a soft (not fixed) state
>> **Eventual consistency:**
The system will eventually become consistent (no more transactions)

Key-Value Database

- A simple hash table accessible only through its primary key
- Basically a table with two columns: ID and VALUE
- The value is a blob that the data store just stores: it can be text, JSON, XML, and anything else.
- Operations:
 - get the value for the key
 - put a value for a key (if the key already exists the corresponding value is overwritten)
 - delete a key from the data store

notable examples: Riak, Redis, Memcached, Berkeley DB
HamsterDB , Amazon DynamoDB, Project Voldemort

Key/Value Stores are usually used for this:
>> Storing Session Information
>> User Profiles, Preferences, Configs
>> Shopping carts (LMAO)
They are **NOT** to be used in these cases:
<< relationship among data , multioperation transactions
<< query by data, operation by sets

NoSQL Aggregation database

> Arrays (Array Store)
> Dictionaries (Key/Value Stores)
> nested Structures(Document-databases)

NoSQL Document Store

>> Maps a key to a structured Document
>> flexible schema , Document stored in JSON or BSON
>> Examples: MongoDB, CouchDB
term comparison to RDBMS
instance -i instance/database, table -t collection
row -i Document , rowid -i _id/objectId

MongoDB
>> JSON, partitioning via sharding, FOSS :)
>> own query language, Document Store (no Schema)

```
graph TD
    DBs[0 or more Databases] --> Collections[0 or more Collections]
    Collections --> Documents[0 or more Documents]
    Documents --> Fields[1 or more Fields]
```

var class = {
 _id: ObjectId ("509980df3"),
 course: {code: "Dbs2",
 title: "Advanced DB"},
 year: 2016,
 students: ["Peter", "Manuel",...],
 no_of_st: 34
}

- The document contains values of varying types
 - The primary key is _id and it is of the ObjectId type,
 - The field course is a subdocument, and
 - students is an array of strings

The documents are actually stored as BSON
the binary version of JSON with a bit more data.
Also there are certain restriction on naming...

_id

- The field _id is reserved for use as a document's primary key
 - Its value must be unique in the collection,
 - It is immutable,
 - May be of any type other than an array or a regular expression type
 - MongoDB creates a unique index on the _id field during the creation of a collection.
 - It is always the first field in the documents
- The following are common options for storing values for _id:
 - Use an ObjectId,
 - Use a natural unique identifier, if available
 - This saves space and avoids an additional index,
 - Generate an auto-incrementing number
- ObjectId** is a 12-byte BSON type, constructed using:
 - A 4-byte value representing the seconds since the Unix epoch,
 - A 3-byte machine identifier,
 - A 2-byte process id, and
 - A 3-byte counter, starting with a random value
- ObjectIds are small, most likely unique, and fast to generate
- MongoDB uses ObjectIds as the default value for the _id field if the _id field is not specified by a client
- Additional benefits of using ObjectIds for the _id field:
 - In the mongo shell, you can access the creation time of the ObjectId, using the getTimestamp() method,
 - Sorting on ObjectId values is roughly equivalent to sorting by creation time.


MongoDB Queries

- Expressed via JSON with simple constructs

```
// In orders
{"orderId":99,
 "customerId":"883c2c5b4e5b",
 "orderDate":"2014-04-23",
 "orderItems":[
  {"product":{"id":27,"name":"NoSQL Distilled"}, "price": 32.45 }
  {"product":{"id":55,"name":"Java 4 all"}, "price": 41.33 }
 ],
}
db.orders.find()
db.orders.find({"customerId":"883c2c5b4e5b"})
db.order.find({"customerId":"883c2c5b4e5b"},
[orderId:1,orderDate:1]) db.orders.find({"orderItems.product.name":"NoSQL/})
```

NoSQL Graph Stores

made for compact objects with typed aggregations
sometimes without schema -> implicit schema
or not uniform data



<p>easiest way to convert PSQL to JSON is a temporary table.</p> <pre>create temporary table angprojj as select ang.persnr, jsonb_build_object('persnr', ang.persnr, 'name', min(ang.name), 'projects', jsonb_agg(trim(proj.bezeichnung))) as angwithproj</pre> <p>joins are omitted, would be below...</p> <p>you can also directly <code>select json.build.object</code> or <code>select json.build.object(jsonb_agg(tmp))</code> also note the <code>jsonb_agg(trim (proj.bezeichnung))</code> for simplicity, the joins on the picture are removed.</p> <p>The <code>jsonb agg</code> is necessary for aggregation otherwise the join wouldn't work for JSON.</p>
<p>JSON-Daten abfragen</p> <pre>select persnr, angwithproj from angprojj where angwithproj->>'persnr' = 1001::text; select persnr, angwithproj from angprojj where (angwithproj->>'name') like 'Marxer%' select persnr, jsonb_pretty(angwithproj) from angprojj where angwithproj->'projects' @> to_jsonb('Uranus'::text) Get as text: select persnr, angwithproj->'projects' as projects from angprojj; => 23 rows CROSS JOIN zweier Tabellen. jsonb_array_elements_text() gibt „setof text“ zurück: select persnr, angwithproj->>'name' as persname, value as projname from angprojj, jsonb_array_elements_text(angwithproj->'projects'); => 29 rows JOIN LATERAL = CROSS JOIN und Boolean: Output identisch mit oben : select persnr, angwithproj->>'name' as persname, value as projname from angprojj cross join lateral jsonb_array_elements_text(angwithproj->'projects'); => 29 rows</pre>
<p>JSON — SQL</p>