

Functions in PGPLSQL
functions in psql create or replace function funcName() return s returntype as \$\$ begin raise notice 'Hello Birb!'; end; \$\$ language langName The two \$ are always necessary. Also note the returns with an s and the language at the end. Which MUST be a PROCEDURAL LANGUAGE , so c++ doesn't work here.
Parameters are handled like in any language func(x bigint, y bigint) you can also define multiple return types func(variadic a numeric[]) or a generic return func(param anyelement)
Variable Declaration: returns void as \$\$ DECLARE x bigint; y bigint; BEGIN
Variable manipulation: x := 6 + 4
if: IF n = 0 THEN RETURN 1; (optional)ELSE RETURN 2; END IF; ELSIF also possible (note elsif not elseif)
case x when 1, 2 then msg := 'one or two'; end case; essentially this checks if x is 1 OR 2
case when x between 0 and 10 then similar but with a range, both can be simulated by if else.
Exceptions: BEGIN z:= x / y; EXCEPTION WHEN division-by-zero THEN z:= 0; (or error rather) END; if you want to catch all: WHEN others THEN often used after exception: RAISE; (show error)
For Loop: For var IN query LOOP statements END LOOP; for r in SELECT * FROM ang LOOP RETURN NEXT r; END LOOP; RETURN; END; note that the return next doesn't return you store it in a buffer and return it at the end of the function. for infinite loops: FOR i IN 1...max LOOP;
update and insert: INSERT INTO ANG VALUES(...); UPDATE ang set salary = salary + 500 where name = 'dashie'; interestingly, after the where name = 'dashie' you can use if not found then (handle error) this allows for easier error handling.
queries: execute 'SELECT * from ang' —— into result; return result; END;
comments are done by either – or /* */ for multiline
anonymous function: you can omit the name and just write do \$\$...
cursor: declare curs CURSOR FOR query; BEGIN OPEN curs; LOOP do something CLOSE curs; END; Cursors are essentially just iterables. cursors can also be unbound curs1 refcursor or they can be parameterized curs3 cursor(arg)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row) arrays: SELECT '1,2,3':int[] or SELECT ARRAY[1,2,3] var int[] only in variable declaration. !!arrays start with 1 in psql !!
return types: all of the above AND void, SETOF type (array of a type), TABLE, Trigger
Arrays: Accessoren create table tictactoe as (select 1 as id, array[['z1 k1', 'z1 k2'], 'z2 k1', 'z2 k2'], 'z3 k1', 'z3 k2']] as board); Index Query: intuitiv wie eine Koordinate ("1-basiert": Start mit 1 nicht 0): select board[1][1] from tictactoe; -- z1 k1 Slice Query: "Untergrenze:Obergrenze" für jede Dimension: select board[2:3][1:1] from tictactoe; -- {{z2 k1},{z3 k1}} Max-Bound-Abkürzung "[2]" vermeiden (Verwechslungsgefahr), besser [1:2] select board[2:3][2] from tictactoe; -- [2]=>[1:2] -- {{z2 k1,z2 k2},{z3 k1,z3 k2}} Suche mit ANY: select * from tictactoe where 'z2 k2' = any (board); -- 1;{{z1 k1,z1 k2},{z2 k1,z2 k2},{z3 k1,z3 k2}}
Arrays: Operatoren • «ls equal»: = – SELECT ARRAY[1,2,3] = ARRAY[1,2,3]; true – SELECT ARRAY[3,2,1] = ARRAY[1,2,3]; false • «Contains»: @> – SELECT ARRAY[1,7,4,2,6] @> ARRAY[2,7]; • «ls contained by»: <@ – SELECT ARRAY[2,7] <@ ARRAY[1,7,4,2,6]; • «Overlaps»: && – SELECT ARRAY[1,4,3] && ARRAY[2,1]
hstore / map: Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL);

stored procedures are nothing but a chaining of functions: Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row) some good to know things: plain SQL is more efficient. write variables lower case for sql use cast over typename -> not select date '2022-06-07'
Triggers Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row) statement is once, row means once per row, aka for the entire table.
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Create und Insert: CREATE TABLE test (id integer, col2 hstore, col3 text); INSERT INTO test VALUES (1, 'a=>123, b=>foo, c=>bar':hstore, NULL); SELECT * FROM test; id col2 col3 ----- 1 "a"=>"123", "b"=>"foo", "c"=>"bar" null (1 row)
Return types: RETURN NEW -> returns a new table/row RETURN OLD -> returns the old table/row (but could change other rows!) RETURN NULL -> cancel operation.
running order: before statement, before row, after row, after statement -> and of course alphabetically.
inside the Trigger functions you can use the variables that don't matter aka can be ANY -> user or the entered user from the trigger -> NEW.user or explicitly the old one -> OLD.user and last user defined stuff like -> SELECT 'I' which just places an I as the variable or something like now() for timestamps.
Triggers make the database slower and harder to maintain. some databases therefore let you disable them if you want. On a table basis.

Optimization and indexing
The basics of indexing is that it saves time on queries, but it uses more space, and needs to be redone on update/insert.
Data stored in Pages / Heap (Collection of Pages) Indexing on either can be possible. ex: Page index -> Primary key index ex: Heap Index -> index on tables PSQL does table cluster indexing instead of integrated indexing -> key value (both indexes)
Hash Index: just like hashmap in programming good for single or small multiple queries bad for entire tables etc -> collision hashing might take a long time with a lot of data.
Bitmap indexing: Bitmap -> 0 1 stores Booleans/Enums very fast read / slow update in postgres only implicit use Brin instead.
Brin indexing: